

```
In [1]: 1 from keras.datasets import mnist
2 from keras.models import Sequential
3 from keras.layers import Dropout
4 from keras.layers import BatchNormalization
5 from keras.layers import Dense
6 from keras.layers import Flatten
7 from keras.layers import Conv2D, MaxPooling2D
8 from keras import backend as K
9 from keras.initializers import he_normal
10
11 import numpy as np
12 import pandas as pd
13 import matplotlib.pyplot as plt
14 import seaborn as sns
15
16 %matplotlib inline
```

Using TensorFlow backend.

```
In [2]: 1 #getting the size of input image
2 img_rows,img_cols = 28,28
3
4 #splitting the dataset
5 (X_train,Y_train),(X_test,Y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<https://s3.amazonaws.com/img-datasets/mnist.npz>)

11493376/11490434 [=====] - 11s 1us/step

```
In [3]: 1 print('Number of images in training set are:',X_train.shape[0],'And shape of each image is {} x {}'.format(X
2 print('Number of images in test set are:',X_test.shape[0],'And shape of each image is {} x {}'.format(X_test
```

Number of images in training set are: 60000 And shape of each image is 28 x 28
Number of images in test set are: 10000 And shape of each image is 28 x 28

```
In [0]: 1 #as keras backend perceives incoming input to be of 4 dimensional tensor that's why we need to reshape the s
2
3 if K.image_data_format() == 'channels_first':
4     X_train = X_train.reshape(X_train.shape[0],1,img_row,img_cols)
5     X_test = X_test.reshape(X_test.shape[0],1,img_rows,img_cols)
6     input_shape = (1,img_rows,img_cols)
7 else:
8     X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
9     X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
10    input_shape = (img_rows, img_cols, 1)
11
```

```
In [5]: 1 X_train.shape
2 X_test.shape
3 #backend percieves other than channels first
```

```
Out[5]: (10000, 28, 28, 1)
```

```
In [6]: 1 #display a plot number
2 plt.figure(figsize=(10,10))
3 rows = 3
4 cols= 3
5 for i in range(rows*cols):
6     plt.subplot(rows,cols,i+1)
7     plt.imshow(X_train[i].reshape(28,28),cmap="afmhot",interpolation="none")
8     print(Y_train[i])
9     #reshaping from 1D to 2D array
10    plt.xticks([])
11    plt.yticks([])
```

5
0
4
1
9
2
1
3
1



```
In [7]: 1 X_train = X_train.astype('float32')
2 X_test = X_test.astype('float32')
3 X_train /= 255
4 X_test /= 255
5 print('Shape of training data is',X_train.shape)
6 print(X_train.shape[0], 'train_samples')
7 print(X_test.shape[0], 'test_samples')
```

Shape of training data is (60000, 28, 28, 1)
60000 train_samples
10000 test_samples

```
In [8]: 1 #converting the labels into vectos using one hot encoding
2
3 from keras.utils import to_categorical
4
5 num_classes = 10
6 Y_train = to_categorical(Y_train,num_classes)
7 Y_test = to_categorical(Y_test,num_classes)
8 Y_train.shape
```

Out[8]: (60000, 10)

```
In [0]: 1 def classifier_1(n_kernel):
2     model = Sequential()
3     model.add(Conv2D(64, kernel_size = (n_kernel, n_kernel), activation = 'relu', kernel_initializer = he_normal(s
4     model.add(Conv2D(32, (n_kernel, n_kernel), activation = 'relu', kernel_initializer = he_normal(seed = None)))
5     model.add(MaxPooling2D(pool_size=(2, 2)))
6     model.add(BatchNormalization())
7     model.add(Dropout(0.3))
8
9     #adding a flattening layer
10    model.add(Flatten())
11    model.add(Dense(128, activation = 'relu', kernel_initializer = he_normal(seed = None)))
12    model.add(BatchNormalization())
13    model.add(Dropout(0.5))
14
15    #addding the output layer
16    model.add(Dense(10, activation = 'softmax'))
17
18    model.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
19
20    return model
21
22
23
```

```
In [0]: 1 import warnings
2     warnings.filterwarnings('ignore')
```

```

In [0]: 1 from keras.wrappers.scikit_learn import KerasClassifier
2 from sklearn.model_selection import GridSearchCV
3 import time
4
5 start_time = time.time()
6
7 batch_size = 100
8 epochs = 20
9 #writing the fucntion to perform parameter tuning and pass the argument of different kernel sizes as a list
10 n_kernel = [2,3,5,7,11]
11
12
13 model_gs = KerasClassifier(build_fn = classifier_1,batch_size = batch_size,epochs = epochs,verbose = 0)
14 param_grid = dict(n_kernel = n_kernel)
15 grid = GridSearchCV(estimator = model_gs,param_grid = param_grid,cv = 5,verbose = 1)
16 grid_result = grid.fit(X_train,Y_train)
17
18 end_time = time.time()
19 print('Time taken for execution of the hypereparamter searching is:',(end_time - start_time))
20
21
22 print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
23 means = grid_result.cv_results_['mean_test_score']
24 stds = grid_result.cv_results_['std_test_score']
25 params = grid_result.cv_results_['params']
26 for mean, stdev, param in zip(means, stds, params):
27     print("%f (%f) with: %r" % (mean, stdev, param))

```

Fitting 5 folds for each of 5 candidates, totalling 25 fits

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4479: The name tf.truncated_normal is deprecated. Please use tf.random.truncated_normal instead.

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.


```

In [11]: 1 #after getting the best kernel_size
          2 # we need to evaluate it on the test data
          3
          4 #best_kernel_size is 7
          5 batch_size = 100
          6 epochs = 20
          7 model = Sequential()
          8 model.add(Conv2D(64,kernel_size = (7,7),activation = 'relu',kernel_initializer = he_normal(seed = None),inp
          9 model.add(Conv2D(32,(7,7),activation = 'relu',kernel_initializer = he_normal(seed = None)))
         10 model.add(MaxPooling2D(pool_size=(2, 2)))
         11 model.add(BatchNormalization())
         12 model.add(Dropout(0.3))
         13
         14 #adding a flattening layer
         15 model.add(Flatten())
         16 model.add(Dense(128,activation = 'relu',kernel_initializer = he_normal(seed = None)))
         17 model.add(BatchNormalization())
         18 model.add(Dropout(0.5))
         19
         20 #adding the output layer
         21 model.add(Dense(10,activation = 'softmax'))
         22 print(model.summary())
         23
         24
         25 model.compile(loss = 'categorical_crossentropy',optimizer = 'adam', metrics = ['accuracy'])
         26 history = model.fit(X_train,Y_train,batch_size = batch_size,epochs = epochs,validation_data = (X_test,Y_test

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4479: The name tf.truncated_normal is deprecated. Please use tf.random.truncated_normal instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4267: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:2041: The name tf.nn.fused_batch_norm is deprecated. Please use tf.compat.v1.nn.fused_batch_norm instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 22, 22, 64)	3200
conv2d_2 (Conv2D)	(None, 16, 16, 32)	100384
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 32)	128
dropout_1 (Dropout)	(None, 8, 8, 32)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_1 (Dense)	(None, 128)	262272
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
Total params: 367,786		
Trainable params: 367,466		
Non-trainable params: 320		

None

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 164us/step - loss: 0.1793 - acc: 0.9458 - val_loss: 0.0477 - val_acc: 0.9845

Epoch 2/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0652 - acc: 0.9811 - val_loss: 0.0315 - val_acc: 0.9908

Epoch 3/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0490 - acc: 0.9851 - val_loss: 0.0253 - val_acc: 0.9921

Epoch 4/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0387 - acc: 0.9881 - val_loss: 0.0268 - val_acc: 0.9923

Epoch 5/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0326 - acc: 0.9899 - val_loss: 0.0230 - val_acc: 0.9921

Epoch 6/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0264 - acc: 0.9917 - val_loss: 0.0231 - val_acc: 0.9927

Epoch 7/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0255 - acc: 0.9918 - val_loss: 0.0227 - val_acc: 0.9917

Epoch 8/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0248 - acc: 0.9922 - val_loss: 0.0207 - val_acc: 0.9931

Epoch 9/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0218 - acc: 0.9929 - val_loss: 0.0214 - val_acc: 0.9930

Epoch 10/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0176 - acc: 0.9944 - val_loss: 0.0195 - val_acc: 0.9938

Epoch 11/20

60000/60000 [=====] - 6s 94us/step - loss: 0.0170 - acc: 0.9944 - val_loss: 0.0211 - val_acc: 0.9931

Epoch 12/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0165 - acc: 0.9943 - val_loss: 0.0217 - val_acc: 0.9928

```
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0154 - acc: 0.9949 - val_loss: 0.0281 - val_acc: 0.9911
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0149 - acc: 0.9950 - val_loss: 0.0181 - val_acc: 0.9947
Epoch 15/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0133 - acc: 0.9960 - val_loss: 0.0197 - val_acc: 0.9934
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0113 - acc: 0.9964 - val_loss: 0.0192 - val_acc: 0.9940
Epoch 17/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0109 - acc: 0.9966 - val_loss: 0.0208 - val_acc: 0.9940
Epoch 18/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0102 - acc: 0.9967 - val_loss: 0.0205 - val_acc: 0.9941
Epoch 19/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0104 - acc: 0.9967 - val_loss: 0.0242 - val_acc: 0.9928
Epoch 20/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0100 - acc: 0.9967 - val_loss: 0.0196 - val_acc: 0.9951
```

In [0]:

1

Though we got an absolutely amazing accuracy of 99.31 % on test data, in the next model that we will build, data augmentation would be used on the data for better understanding the data

```
In [0]: 1 import pickle
        2
        3 import pickle
        4 def savetofile(obj,filename):
        5     pickle.dump(obj,open(filename+".p",'wb'))
        6
        7 def openfromfile(filename):
        8     temp = pickle.load(open(filename+".p",'rb'))
        9     return temp
       10
```

```
In [0]: 1 history1 = savetofile(history,'history_1')
        2 model1 = savetofile(model,'model1')
```

```
In [14]: 1 model1 = openfromfile('model1')
        2 score_model1 = model1.evaluate(X_test,Y_test)
        3 print('Loss on test data is',score_model1[0])
        4 print('Accuracy on test data is',score_model1[1])
```

```
10000/10000 [=====] - 1s 89us/step
Loss on test data is 0.01964392685413427
Accuracy on test data is 0.9951
```

```
In [15]: 1 X_train.shape
```

```
Out[15]: (60000, 28, 28, 1)
```

In [16]:

```
1  #CNN with 5 Layers
2
3  model = Sequential()
4
5  #1st and 2nd layer with maxpooling layer
6  model.add(Conv2D(32,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None),input_sh
7  model.add(Conv2D(32,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None)))
8  model.add(MaxPooling2D(2,2))
9  model.add(BatchNormalization())
10 model.add(Dropout(0.5))
11
12 #3rd and 4th layer with maxpooling layer
13 model.add(Conv2D(64,kernel_size = 2,activation = 'relu',kernel_initializer = he_normal(seed = None)))
14 model.add(Conv2D(64,kernel_size = 2,activation = 'relu',kernel_initializer = he_normal(seed = None)))
15 model.add(MaxPooling2D(2,2))
16 model.add(BatchNormalization())
17 model.add(Dropout(0.5))
18
19
20 #5th Layer
21 model.add(Conv2D(128,kernel_size = 5,activation = 'relu',kernel_initializer = he_normal(seed = None)))
22 model.add(BatchNormalization())
23 model.add(Dropout(0.4))
24
25 #flattening and adding another 2 dense layers with relu activation
26 model.add(Flatten())
27 model.add(Dense(256,activation = 'relu',kernel_initializer = he_normal(seed = None)))
28 model.add(Dropout(0.4))
29
30
31 model.add(Dense(128,activation = 'relu',kernel_initializer = he_normal(seed = None)))
32 model.add(Dropout(0.4))
33
34
35 #output layer
36 model.add(Dense(10,activation = 'softmax'))
37
38 print(model.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
conv2d_4 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 32)	0
batch_normalization_3 (Batch Normalization)	(None, 12, 12, 32)	128
dropout_3 (Dropout)	(None, 12, 12, 32)	0
conv2d_5 (Conv2D)	(None, 11, 11, 64)	8256
conv2d_6 (Conv2D)	(None, 10, 10, 64)	16448
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 64)	0
batch_normalization_4 (Batch Normalization)	(None, 5, 5, 64)	256
dropout_4 (Dropout)	(None, 5, 5, 64)	0
conv2d_7 (Conv2D)	(None, 1, 1, 128)	204928
batch_normalization_5 (Batch Normalization)	(None, 1, 1, 128)	512
dropout_5 (Dropout)	(None, 1, 1, 128)	0
flatten_2 (Flatten)	(None, 128)	0
dense_3 (Dense)	(None, 256)	33024
dropout_6 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 128)	32896
dropout_7 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 10)	1290
Total params: 307,306		
Trainable params: 306,858		
Non-trainable params: 448		

None

In [17]:

```

1 #optimization for compilation
2
3 model.compile(loss = 'categorical_crossentropy',optimizer = 'adam',metrics = ['accuracy'])
4 history = model.fit(X_train,Y_train,batch_size = batch_size,epochs = epochs,validation_data = (X_test,Y_test)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.7325 - acc: 0.7727 - val_loss: 0.1535
- val_acc: 0.9553

Epoch 2/20

60000/60000 [=====] - 6s 101us/step - loss: 0.1859 - acc: 0.9471 - val_loss: 0.1055
- val_acc: 0.9669

Epoch 3/20

60000/60000 [=====] - 6s 102us/step - loss: 0.1418 - acc: 0.9615 - val_loss: 0.0475
- val_acc: 0.9845

Epoch 4/20

60000/60000 [=====] - 6s 101us/step - loss: 0.1171 - acc: 0.9675 - val_loss: 0.0447
- val_acc: 0.9873

Epoch 5/20

60000/60000 [=====] - 6s 102us/step - loss: 0.1005 - acc: 0.9726 - val_loss: 0.0389
- val_acc: 0.9886

Epoch 6/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0924 - acc: 0.9747 - val_loss: 0.0388
- val_acc: 0.9886

Epoch 7/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0879 - acc: 0.9760 - val_loss: 0.0454
- val_acc: 0.9877

Epoch 8/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0760 - acc: 0.9791 - val_loss: 0.0380
- val_acc: 0.9880

Epoch 9/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0752 - acc: 0.9788 - val_loss: 0.0360
- val_acc: 0.9899

Epoch 10/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0696 - acc: 0.9808 - val_loss: 0.0314
- val_acc: 0.9915

Epoch 11/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0678 - acc: 0.9821 - val_loss: 0.0278
- val_acc: 0.9922

Epoch 12/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0632 - acc: 0.9829 - val_loss: 0.0326
- val_acc: 0.9917

```

Epoch 13/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0591 - acc: 0.9837 - val_loss: 0.0283
- val_acc: 0.9931
Epoch 14/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0564 - acc: 0.9841 - val_loss: 0.0291
- val_acc: 0.9916
Epoch 15/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0541 - acc: 0.9846 - val_loss: 0.0304
- val_acc: 0.9922
Epoch 16/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0562 - acc: 0.9846 - val_loss: 0.0263
- val_acc: 0.9929
Epoch 17/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0488 - acc: 0.9862 - val_loss: 0.0262
- val_acc: 0.9917
Epoch 18/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0497 - acc: 0.9863 - val_loss: 0.0242
- val_acc: 0.9936
Epoch 19/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0504 - acc: 0.9861 - val_loss: 0.0225
- val_acc: 0.9932
Epoch 20/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0490 - acc: 0.9867 - val_loss: 0.0232
- val_acc: 0.9937

```

```

In [0]: 1 history2 = savetofile(history, 'history2')
        2 model2 = savetofile(model, 'model2')

```

```

In [19]: 1 model2 = openfromfile('model2')
        2 score_model2 = model2.evaluate(X_test, Y_test)
        3 print('Loss on test data is', score_model2[0])
        4 print('Accuracy on test data is', score_model2[1])

```

```

10000/10000 [=====] - 1s 112us/step
Loss on test data is 0.023159436680417637
Accuracy on test data is 0.9937

```

In [20]:

```
1 #arch 3
2 #building a 7 layer of Conv2d ,7 layers of maxpooling and 2 dense layers
3
4 model = Sequential()
5
6 model.add(Conv2D(16,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None),input_sh
7 model.add(MaxPooling2D(1,1))
8 model.add(BatchNormalization())
9 model.add(Dropout(0.5))
10
11
12 model.add(Conv2D(32,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None)))
13 model.add(MaxPooling2D(1,1))
14 model.add(BatchNormalization())
15 model.add(Dropout(0.5))
16
17
18 model.add(Conv2D(64,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None)))
19 model.add(MaxPooling2D(1,1))
20 model.add(BatchNormalization())
21 model.add(Dropout(0.5))
22
23
24 model.add(Conv2D(96,kernel_size = 5,activation = 'relu',kernel_initializer = he_normal(seed = None)))
25 model.add(MaxPooling2D(1,1))
26 model.add(BatchNormalization())
27 model.add(Dropout(0.3))
28
29
30
31
32 model.add(Conv2D(128,kernel_size = 5,activation = 'relu',kernel_initializer = he_normal(seed = None)))
33 model.add(MaxPooling2D(2,2))
34 model.add(BatchNormalization())
35 model.add(Dropout(0.3))
36
37
38 model.add(Conv2D(160,kernel_size = 5,activation = 'relu',kernel_initializer = he_normal(seed = None)))
39 model.add(MaxPooling2D(2,2))
40 model.add(BatchNormalization())
41 model.add(Dropout(0.3))
42
```

```

43
44
45 #flattening and adding another 2 dense layers with relu activation
46 model.add(Flatten())
47 model.add(Dense(256,activation = 'relu',kernel_initializer = he_normal(seed = None)))
48 model.add(BatchNormalization())
49 model.add(Dropout(0.3))
50
51
52 model.add(Dense(128,activation = 'relu',kernel_initializer = he_normal(seed = None)))
53 model.add(Dropout(0.3))
54
55
56 #output layer
57 model.add(Dense(10,activation = 'softmax'))
58
59 print(model.summary())

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_4 (MaxPooling2)	(None, 26, 26, 16)	0
batch_normalization_6 (Batch Normalization)	(None, 26, 26, 16)	64
dropout_8 (Dropout)	(None, 26, 26, 16)	0
conv2d_9 (Conv2D)	(None, 24, 24, 32)	4640
max_pooling2d_5 (MaxPooling2)	(None, 24, 24, 32)	0
batch_normalization_7 (Batch Normalization)	(None, 24, 24, 32)	128
dropout_9 (Dropout)	(None, 24, 24, 32)	0
conv2d_10 (Conv2D)	(None, 22, 22, 64)	18496
max_pooling2d_6 (MaxPooling2)	(None, 22, 22, 64)	0

batch_normalization_8 (Batch Normalization)	(None, 22, 22, 64)	256
dropout_10 (Dropout)	(None, 22, 22, 64)	0
conv2d_11 (Conv2D)	(None, 18, 18, 96)	153696
max_pooling2d_7 (MaxPooling2D)	(None, 18, 18, 96)	0
batch_normalization_9 (Batch Normalization)	(None, 18, 18, 96)	384
dropout_11 (Dropout)	(None, 18, 18, 96)	0
conv2d_12 (Conv2D)	(None, 14, 14, 128)	307328
max_pooling2d_8 (MaxPooling2D)	(None, 7, 7, 128)	0
batch_normalization_10 (Batch Normalization)	(None, 7, 7, 128)	512
dropout_12 (Dropout)	(None, 7, 7, 128)	0
conv2d_13 (Conv2D)	(None, 3, 3, 160)	512160
max_pooling2d_9 (MaxPooling2D)	(None, 1, 1, 160)	0
batch_normalization_11 (Batch Normalization)	(None, 1, 1, 160)	640
dropout_13 (Dropout)	(None, 1, 1, 160)	0
flatten_3 (Flatten)	(None, 160)	0
dense_6 (Dense)	(None, 256)	41216
batch_normalization_12 (Batch Normalization)	(None, 256)	1024
dropout_14 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 128)	32896
dropout_15 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 1,074,890		

Trainable params: 1,073,386

Non-trainable params: 1,504

None

In [21]:

```

1 #optimization for compilation
2
3 model.compile(loss = 'categorical_crossentropy',optimizer = 'adam',metrics = ['accuracy'])
4 history = model.fit(X_train,Y_train,batch_size = batch_size,epochs = epochs,validation_data = (X_test,Y_test)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 22s 369us/step - loss: 0.5196 - acc: 0.8393 - val_loss: 0.097

4 - val_acc: 0.9700

Epoch 2/20

60000/60000 [=====] - 20s 329us/step - loss: 0.1517 - acc: 0.9543 - val_loss: 0.060

7 - val_acc: 0.9809

Epoch 3/20

60000/60000 [=====] - 19s 322us/step - loss: 0.1020 - acc: 0.9699 - val_loss: 0.046

4 - val_acc: 0.9860

Epoch 4/20

60000/60000 [=====] - 19s 321us/step - loss: 0.0813 - acc: 0.9758 - val_loss: 0.037

7 - val_acc: 0.9884

Epoch 5/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0729 - acc: 0.9784 - val_loss: 0.031

7 - val_acc: 0.9902

Epoch 6/20

60000/60000 [=====] - 19s 322us/step - loss: 0.0626 - acc: 0.9813 - val_loss: 0.036

6 - val_acc: 0.9888

Epoch 7/20

60000/60000 [=====] - 19s 322us/step - loss: 0.0561 - acc: 0.9832 - val_loss: 0.022

2 - val_acc: 0.9933

Epoch 8/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0527 - acc: 0.9846 - val_loss: 0.036

6 - val_acc: 0.9906

Epoch 9/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0464 - acc: 0.9862 - val_loss: 0.021

4 - val_acc: 0.9942

Epoch 10/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0458 - acc: 0.9866 - val_loss: 0.025

4 - val_acc: 0.9932

Epoch 11/20

60000/60000 [=====] - 19s 319us/step - loss: 0.0414 - acc: 0.9881 - val_loss: 0.024

3 - val_acc: 0.9928

Epoch 12/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0386 - acc: 0.9885 - val_loss: 0.029

2 - val_acc: 0.9919

```

Epoch 13/20
60000/60000 [=====] - 19s 320us/step - loss: 0.0354 - acc: 0.9894 - val_loss: 0.021
6 - val_acc: 0.9934
Epoch 14/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0350 - acc: 0.9893 - val_loss: 0.023
5 - val_acc: 0.9933
Epoch 15/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0325 - acc: 0.9905 - val_loss: 0.026
3 - val_acc: 0.9929
Epoch 16/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0290 - acc: 0.9911 - val_loss: 0.021
6 - val_acc: 0.9937
Epoch 17/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0290 - acc: 0.9910 - val_loss: 0.021
1 - val_acc: 0.9940
Epoch 18/20
60000/60000 [=====] - 19s 321us/step - loss: 0.0277 - acc: 0.9914 - val_loss: 0.023
9 - val_acc: 0.9931
Epoch 19/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0245 - acc: 0.9923 - val_loss: 0.019
0 - val_acc: 0.9940
Epoch 20/20
60000/60000 [=====] - 19s 319us/step - loss: 0.0268 - acc: 0.9919 - val_loss: 0.022
7 - val_acc: 0.9938

```

```

In [0]: 1 history3 = savetofile(history, 'history3')
        2 model3 = savetofile(model, 'model3')

```

```

In [23]: 1 model3 = openfromfile('model3')
        2
        3 score_model3 = model1.evaluate(X_test, Y_test)
        4 print('Loss on test data is', score_model3[0])
        5 print('Accuracy on test data is', score_model3[1])

```

```

10000/10000 [=====] - 1s 92us/step
Loss on test data is 0.01964392685413427
Accuracy on test data is 0.9951


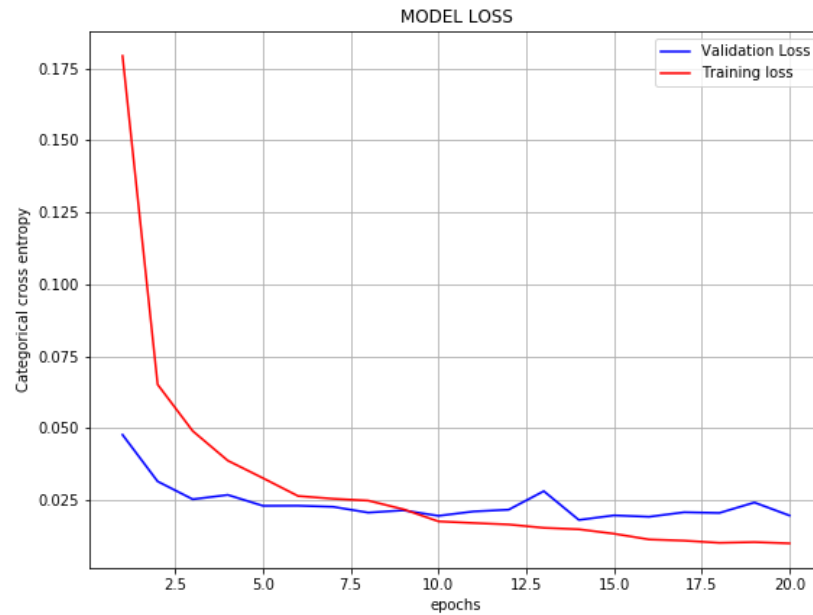
```



```
In [0]: 1 #plotting all the results
2 def plots(model_h):
3
4
5     plt.figure(figsize = (20,15))
6     #plt.grid()
7     x = list(range(1,epochs+1))
8
9     #model Loss
10    plt.subplot(2,2,1)
11    plt.title('MODEL LOSS')
12    plt.grid()
13    plt.plot(x,model_h.history['val_loss'],color = 'b',label = 'Validation Loss')
14    plt.plot(x,model_h.history['loss'],color = 'r',label = 'Training loss')
15    plt.xlabel('epochs')
16    plt.ylabel('Categorical cross entropy')
17    plt.legend()
18
19    #model accuracy
20    plt.subplot(2,2,2)
21    plt.title('MODEL ACCURACY')
22    plt.grid()
23    plt.plot(x,model_h.history['acc'],color = 'b',label = 'Training accuracy')
24    plt.plot(x,model_h.history['val_acc'],color = 'r',label = 'Validation accuracy')
25    plt.xlabel('epochs')
26    plt.ylabel('Accuracy')
27    plt.legend()
28    plt.show()
29
30
31
32
```

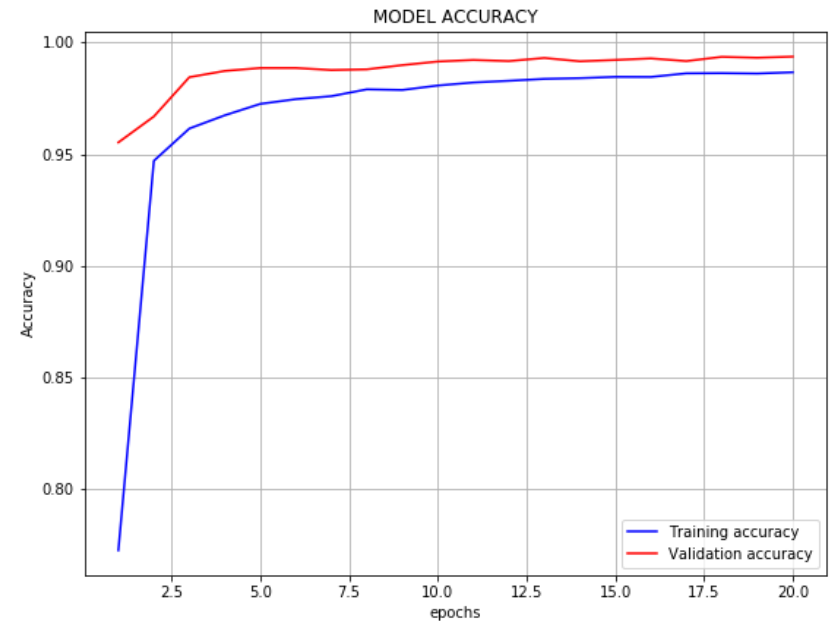

```
In [25]: 1 print('MODEL 1')
          2 history1 = openfromfile('history_1')
          3 plots(history1)
```

MODEL 1





```
In [26]: 1 history2 = openfromfile('history2')
2 print('MODEL 2')
3 plots(history2)
```

MODEL 2



```
In [27]: 1 history3 = openfromfile('history3')
2         print('MODEL 3')
3         plots(history3)
```

MODEL 3



Best Model with Different parameters

- In this section we will use 3 different models each one following the same architecture but with different tuned parameters

MODEL 1: Removing dropout and batchNormalization layers

```
In [28]: 1 model = Sequential()
2
3 #1st and 2nd layer with maxpooling layer
4 model.add(Conv2D(32, kernel_size = 3, activation = 'relu', kernel_initializer = he_normal(seed = None), input_shape=(28, 28, 3)))
5 model.add(Conv2D(32, kernel_size = 3, activation = 'relu', kernel_initializer = he_normal(seed = None)))
6 model.add(MaxPooling2D(2, 2))
7
8
9 #3rd and 4th layer with maxpooling layer
10 model.add(Conv2D(64, kernel_size = 2, activation = 'relu', kernel_initializer = he_normal(seed = None)))
11 model.add(Conv2D(64, kernel_size = 2, activation = 'relu', kernel_initializer = he_normal(seed = None)))
12 model.add(MaxPooling2D(2, 2))
13
14
15
16 #5th layer
17 model.add(Conv2D(128, kernel_size = 5, activation = 'relu', kernel_initializer = he_normal(seed = None)))
18
19
20 #flattening and adding another 2 dense layers with relu activation
21 model.add(Flatten())
22 model.add(Dense(256, activation = 'relu', kernel_initializer = he_normal(seed = None)))
23
24
25 model.add(Dense(128, activation = 'relu', kernel_initializer = he_normal(seed = None)))
26
27
28 #output layer
29 model.add(Dense(10, activation = 'softmax'))
30
31 print(model.summary())
32
33 batch_size = 100
34 epochs = 20
35
36 model.compile(loss = 'categorical_crossentropy', optimizer = 'adam', metrics = ['accuracy'])
37
38 history = model.fit(X_train, Y_train, batch_size = batch_size, epochs = epochs, validation_data = (X_test, Y_test))
39
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_14 (Conv2D)	(None, 26, 26, 32)	320
conv2d_15 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_10 (MaxPooling)	(None, 12, 12, 32)	0
conv2d_16 (Conv2D)	(None, 11, 11, 64)	8256
conv2d_17 (Conv2D)	(None, 10, 10, 64)	16448
max_pooling2d_11 (MaxPooling)	(None, 5, 5, 64)	0
conv2d_18 (Conv2D)	(None, 1, 1, 128)	204928
flatten_4 (Flatten)	(None, 128)	0
dense_9 (Dense)	(None, 256)	33024
dense_10 (Dense)	(None, 128)	32896
dense_11 (Dense)	(None, 10)	1290
Total params: 306,410		
Trainable params: 306,410		
Non-trainable params: 0		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 134us/step - loss: 0.1402 - acc: 0.9563 - val_loss: 0.0402
- val_acc: 0.9872

Epoch 2/20

60000/60000 [=====] - 5s 88us/step - loss: 0.0429 - acc: 0.9866 - val_loss: 0.0407
- val_acc: 0.9864

Epoch 3/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0314 - acc: 0.9903 - val_loss: 0.0309
- val_acc: 0.9906

Epoch 4/20

60000/60000 [=====] - 5s 87us/step - loss: 0.0252 - acc: 0.9918 - val_loss: 0.0289
- val_acc: 0.9910

```
Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0193 - acc: 0.9943 - val_loss: 0.0281
- val_acc: 0.9916
Epoch 6/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0162 - acc: 0.9948 - val_loss: 0.0375
- val_acc: 0.9890
Epoch 7/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0141 - acc: 0.9957 - val_loss: 0.0271
- val_acc: 0.9914
Epoch 8/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0135 - acc: 0.9955 - val_loss: 0.0368
- val_acc: 0.9908
Epoch 9/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0108 - acc: 0.9963 - val_loss: 0.0273
- val_acc: 0.9919
Epoch 10/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0106 - acc: 0.9966 - val_loss: 0.0298
- val_acc: 0.9919
Epoch 11/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0092 - acc: 0.9972 - val_loss: 0.0363
- val_acc: 0.9895
Epoch 12/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0101 - acc: 0.9968 - val_loss: 0.0360
- val_acc: 0.9910
Epoch 13/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0080 - acc: 0.9975 - val_loss: 0.0318
- val_acc: 0.9914
Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0067 - acc: 0.9981 - val_loss: 0.0353
- val_acc: 0.9911
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0066 - acc: 0.9980 - val_loss: 0.0276
- val_acc: 0.9927
Epoch 16/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0091 - acc: 0.9974 - val_loss: 0.0300
- val_acc: 0.9930
Epoch 17/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0053 - acc: 0.9984 - val_loss: 0.0463
- val_acc: 0.9887
Epoch 18/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0060 - acc: 0.9982 - val_loss: 0.0455
- val_acc: 0.9909
Epoch 19/20
```



```
60000/60000 [=====] - 5s 85us/step - loss: 0.0071 - acc: 0.9980 - val_loss: 0.0380  
- val_acc: 0.9922  
Epoch 20/20  
60000/60000 [=====] - 5s 85us/step - loss: 0.0045 - acc: 0.9987 - val_loss: 0.0346  
- val_acc: 0.9927
```

```
In [0]: 1 history_11 = savetofile(history, 'history11')  
2 model_11 = savetofile(model, 'model_11')
```

MODEL2: Using Tanh activation function along with glorot normal weight initialization

```
In [30]: 1 #CNN with 5 Layers
2 from keras.initializers import glorot_normal
3
4 model = Sequential()
5
6 #1st and 2nd Layer with maxpooling Layer
7 model.add(Conv2D(32,kernel_size = 3,activation = 'tanh',kernel_initializer = glorot_normal(seed = None),input_shape=(28,28,3)))
8 model.add(Conv2D(32,kernel_size = 3,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
9 model.add(MaxPooling2D(2,2))
10 model.add(BatchNormalization())
11 model.add(Dropout(0.5))
12
13 #3rd and 4th Layer with maxpooling Layer
14 model.add(Conv2D(64,kernel_size = 2,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
15 model.add(Conv2D(64,kernel_size = 2,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
16 model.add(MaxPooling2D(2,2))
17 model.add(BatchNormalization())
18 model.add(Dropout(0.5))
19
20
21 #5th Layer
22 model.add(Conv2D(128,kernel_size = 5,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
23 model.add(BatchNormalization())
24 model.add(Dropout(0.4))
25
26 #flattening and adding another 2 dense layers with relu activation
27 model.add(Flatten())
28 model.add(Dense(256,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
29 model.add(Dropout(0.4))
30
31
32 model.add(Dense(128,activation = 'tanh',kernel_initializer = glorot_normal(seed = None)))
33 model.add(Dropout(0.4))
34
35
36 #output Layer
37 model.add(Dense(10,activation = 'softmax'))
38
39 print(model.summary())
40
41
42 model.compile(loss = 'categorical_crossentropy',optimizer = 'adam',metrics = ['accuracy'])
```

```

43
44 history = model.fit(X_train,Y_train,batch_size = batch_size,epochs = epochs,validation_data = [X_test,Y_test]
45
46
47
48

```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
=====		
conv2d_19 (Conv2D)	(None, 26, 26, 32)	320
conv2d_20 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_12 (MaxPooling)	(None, 12, 12, 32)	0
batch_normalization_13 (Batch Normalization)	(None, 12, 12, 32)	128
dropout_16 (Dropout)	(None, 12, 12, 32)	0
conv2d_21 (Conv2D)	(None, 11, 11, 64)	8256
conv2d_22 (Conv2D)	(None, 10, 10, 64)	16448
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
batch_normalization_14 (Batch Normalization)	(None, 5, 5, 64)	256
dropout_17 (Dropout)	(None, 5, 5, 64)	0
conv2d_23 (Conv2D)	(None, 1, 1, 128)	204928
batch_normalization_15 (Batch Normalization)	(None, 1, 1, 128)	512
dropout_18 (Dropout)	(None, 1, 1, 128)	0
flatten_5 (Flatten)	(None, 128)	0
dense_12 (Dense)	(None, 256)	33024
dropout_19 (Dropout)	(None, 256)	0

dense_13 (Dense)	(None, 128)	32896
dropout_20 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
=====		
Total params: 307,306		
Trainable params: 306,858		
Non-trainable params: 448		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 10s 159us/step - loss: 0.3995 - acc: 0.8748 - val_loss: 0.2793 - val_acc: 0.9227

Epoch 2/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1660 - acc: 0.9508 - val_loss: 0.1027 - val_acc: 0.9703

Epoch 3/20

60000/60000 [=====] - 6s 103us/step - loss: 0.1301 - acc: 0.9621 - val_loss: 0.0946 - val_acc: 0.9727

Epoch 4/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1192 - acc: 0.9655 - val_loss: 0.4779 - val_acc: 0.8763

Epoch 5/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1069 - acc: 0.9686 - val_loss: 0.1466 - val_acc: 0.9570

Epoch 6/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0999 - acc: 0.9712 - val_loss: 0.0786 - val_acc: 0.9771

Epoch 7/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0942 - acc: 0.9728 - val_loss: 0.0824 - val_acc: 0.9756

Epoch 8/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0954 - acc: 0.9719 - val_loss: 0.1230 - val_acc: 0.9637

Epoch 9/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0918 - acc: 0.9738 - val_loss: 0.1418 - val_acc: 0.9592

Epoch 10/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0879 - acc: 0.9741 - val_loss: 0.0620 - val_acc: 0.9812

```

Epoch 11/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0888 - acc: 0.9744 - val_loss: 0.0724 -
val_acc: 0.9781
Epoch 12/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0839 - acc: 0.9761 - val_loss: 0.0610 -
val_acc: 0.9833
Epoch 13/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0841 - acc: 0.9759 - val_loss: 0.0572 -
val_acc: 0.9850
Epoch 14/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0838 - acc: 0.9760 - val_loss: 0.1263 -
val_acc: 0.9610
Epoch 15/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0807 - acc: 0.9770 - val_loss: 0.0864 -
val_acc: 0.9741
Epoch 16/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0776 - acc: 0.9777 - val_loss: 0.0649 -
val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0778 - acc: 0.9771 - val_loss: 0.0335 -
val_acc: 0.9890
Epoch 18/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0774 - acc: 0.9781 - val_loss: 0.0395 -
val_acc: 0.9875
Epoch 19/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0783 - acc: 0.9780 - val_loss: 0.0443 -
val_acc: 0.9879
Epoch 20/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0761 - acc: 0.9783 - val_loss: 0.0354 -
val_acc: 0.9903

```

```

In [31]: 1 score = model.evaluate(X_test,Y_test)
          2 print('Loss on test data is',score[0])
          3 print('Accuracy on test data is',score[1])

```

```

10000/10000 [=====] - 1s 77us/step
Loss on test data is 0.03541449715318158
Accuracy on test data is 0.9903

```

```

In [0]: 1 history_22 = savetofile(history,'history_22')
          2 model_22 = savetofile(model,'model_22')

```

MODEL 3: Using RMSPROP as optimizer

```
In [33]: 1 from keras.optimizers import RMSprop
2
3 #CNN with 5 layers
4
5 model = Sequential()
6
7 #1st and 2nd layer with maxpooling layer
8 model.add(Conv2D(32,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None),input_shape=(28,28,3)))
9 model.add(Conv2D(32,kernel_size = 3,activation = 'relu',kernel_initializer = he_normal(seed = None)))
10 model.add(MaxPooling2D(2,2))
11 model.add(BatchNormalization())
12 model.add(Dropout(0.5))
13
14 #3rd and 4th layer with maxpooling layer
15 model.add(Conv2D(64,kernel_size = 2,activation = 'relu',kernel_initializer = he_normal(seed = None)))
16 model.add(Conv2D(64,kernel_size = 2,activation = 'relu',kernel_initializer = he_normal(seed = None)))
17 model.add(MaxPooling2D(2,2))
18 model.add(BatchNormalization())
19 model.add(Dropout(0.5))
20
21
22 #5th layer
23 model.add(Conv2D(128,kernel_size = 5,activation = 'relu',kernel_initializer = he_normal(seed = None)))
24 model.add(BatchNormalization())
25 model.add(Dropout(0.4))
26
27 #flattening and adding another 2 dense layers with relu activation
28 model.add(Flatten())
29 model.add(Dense(256,activation = 'relu',kernel_initializer = he_normal(seed = None)))
30 model.add(Dropout(0.4))
31
32
33 model.add(Dense(128,activation = 'relu',kernel_initializer = he_normal(seed = None)))
34 model.add(Dropout(0.4))
35
36
37 #output layer
38 model.add(Dense(10,activation = 'softmax'))
39
40 print(model.summary())
41
42 model.compile(loss = 'categorical_crossentropy',optimizer = 'RMSProp',metrics = ['accuracy'])
```

```

43
44 history = model.fit(X_train,Y_train,batch_size = batch_size,epochs = epochs,validation_data = [X_test,Y_test]
45

```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_24 (Conv2D)	(None, 26, 26, 32)	320
conv2d_25 (Conv2D)	(None, 24, 24, 32)	9248
max_pooling2d_14 (MaxPooling)	(None, 12, 12, 32)	0
batch_normalization_16 (Batch Normalization)	(None, 12, 12, 32)	128
dropout_21 (Dropout)	(None, 12, 12, 32)	0
conv2d_26 (Conv2D)	(None, 11, 11, 64)	8256
conv2d_27 (Conv2D)	(None, 10, 10, 64)	16448
max_pooling2d_15 (MaxPooling)	(None, 5, 5, 64)	0
batch_normalization_17 (Batch Normalization)	(None, 5, 5, 64)	256
dropout_22 (Dropout)	(None, 5, 5, 64)	0
conv2d_28 (Conv2D)	(None, 1, 1, 128)	204928
batch_normalization_18 (Batch Normalization)	(None, 1, 1, 128)	512
dropout_23 (Dropout)	(None, 1, 1, 128)	0
flatten_6 (Flatten)	(None, 128)	0
dense_15 (Dense)	(None, 256)	33024
dropout_24 (Dropout)	(None, 256)	0
dense_16 (Dense)	(None, 128)	32896

dropout_25 (Dropout)	(None, 128)	0
----------------------	-------------	---

dense_17 (Dense)	(None, 10)	1290
------------------	------------	------

=====

Total params: 307,306

Trainable params: 306,858

Non-trainable params: 448

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 157us/step - loss: 0.5762 - acc: 0.8275 - val_loss: 0.1235 - val_acc: 0.9669

Epoch 2/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1622 - acc: 0.9568 - val_loss: 0.0620 - val_acc: 0.9814

Epoch 3/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1273 - acc: 0.9667 - val_loss: 0.0445 - val_acc: 0.9880

Epoch 4/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1194 - acc: 0.9708 - val_loss: 0.0449 - val_acc: 0.9869

Epoch 5/20

60000/60000 [=====] - 6s 104us/step - loss: 0.1043 - acc: 0.9740 - val_loss: 0.0541 - val_acc: 0.9881

Epoch 6/20

60000/60000 [=====] - 6s 103us/step - loss: 0.1003 - acc: 0.9760 - val_loss: 0.0368 - val_acc: 0.9913

Epoch 7/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1030 - acc: 0.9761 - val_loss: 0.0338 - val_acc: 0.9910

Epoch 8/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0971 - acc: 0.9776 - val_loss: 0.0399 - val_acc: 0.9903

Epoch 9/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0977 - acc: 0.9786 - val_loss: 0.0408 - val_acc: 0.9903

Epoch 10/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0964 - acc: 0.9787 - val_loss: 0.0398 - val_acc: 0.9912

Epoch 11/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0967 - acc: 0.9796 - val_loss: 0.0368 - val_acc: 0.9927

```

Epoch 12/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0964 - acc: 0.9797 - val_loss: 0.0406 -
val_acc: 0.9915
Epoch 13/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0951 - acc: 0.9804 - val_loss: 0.0447 -
val_acc: 0.9907
Epoch 14/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0940 - acc: 0.9805 - val_loss: 0.0401 -
val_acc: 0.9911
Epoch 15/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0939 - acc: 0.9803 - val_loss: 0.0366 -
val_acc: 0.9921
Epoch 16/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0875 - acc: 0.9820 - val_loss: 0.0363 - v
al_acc: 0.9923
Epoch 17/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0978 - acc: 0.9809 - val_loss: 0.0464 -
val_acc: 0.9903
Epoch 18/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0975 - acc: 0.9815 - val_loss: 0.0497 -
val_acc: 0.9909
Epoch 19/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0921 - acc: 0.9816 - val_loss: 0.0516 -
val_acc: 0.9894
Epoch 20/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0919 - acc: 0.9821 - val_loss: 0.0466 -
val_acc: 0.9907

```

In [34]:

```

1 score = model.evaluate(X_test,Y_test)
2 print('loss on test data is: ',score[0])
3 print('Accuracy on test data is:',score[1])
4
5 history_33 = savetofile(history,'history_33')
6 model_33 = savetofile(model,'model33')

```

```

10000/10000 [=====] - 1s 78us/step
loss on test data is: 0.046560718773310396
Accuracy on test data is: 0.9907

```

```


In [0]: 1 def diff_ops(model_h,text):
2
3     x = list(range(1,(epochs)+1))#definig the bin size on the x axis as we want to plot for each of the epochs
4     plt.figure(figsize = (18,6))
5     plt.grid()
6     plt.plot(x,model_h.history['val_loss'],'b',label = 'Test loss')
7     plt.plot(x,model_h.history['loss'],'r',label = 'Training Loss')
8     plt.xlabel('Epochs')
9     plt.ylabel('Categorical Cross Entropy')
10    plt.title(text)
11    plt.legend(loc = 'best')
12    plt.show()

```

```


In [36]: 1 text = 'Model_2 Without BatchNormalization and Dropout'
2         diff_ops(openfromfile('history11'),text)
3

```



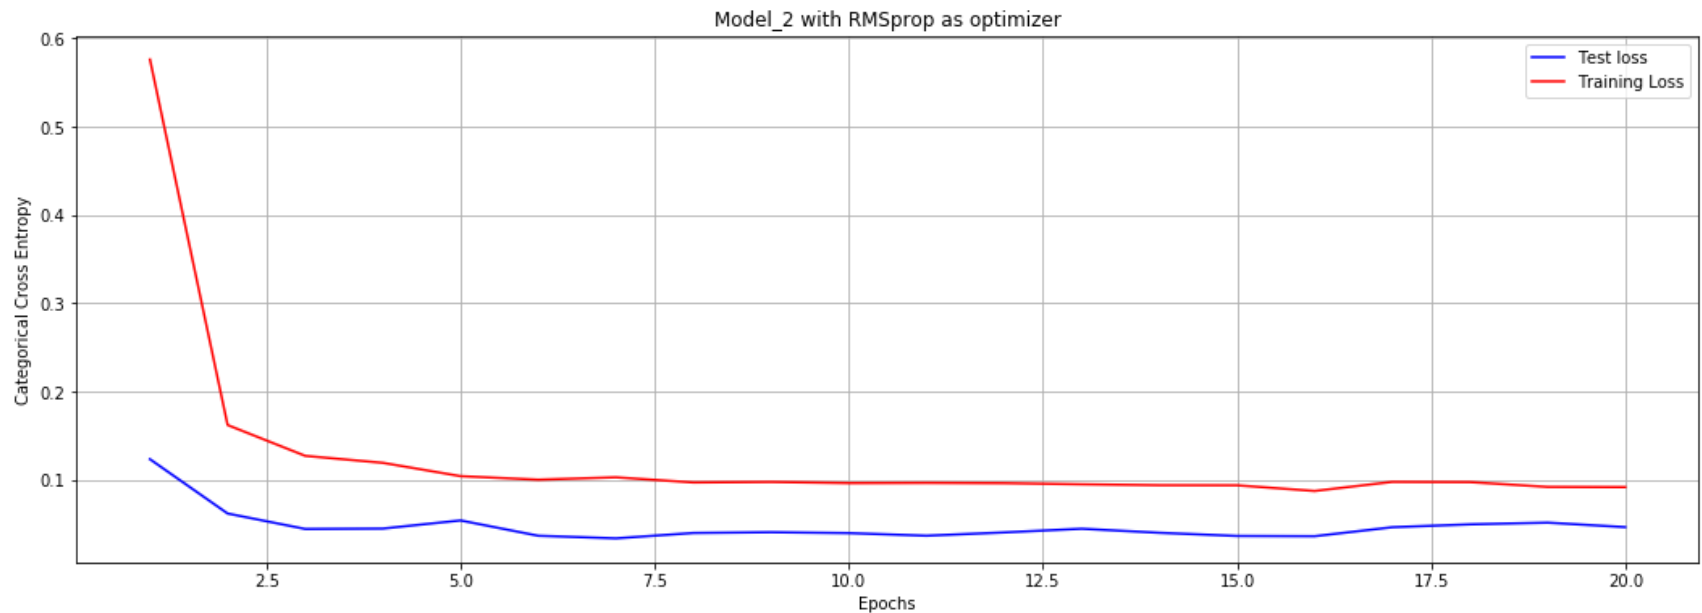
We removed the BatchNormalization layer and the Dropout layers for the best performing model,so as to experiment with overfitting of the model.As it can be clearly seen that model leraned well on the training data ,but performing generalization by this model was not as easy task.Though it gives an accuracy of 99.13 % on test data becuae of low complexity and unsophisticated MNIST data,but it can be clearly observed that the loss is fluctuating with each epoch like an unsmoothed sinusoidal wave.

```
In [37]: 1 #plot_2
2 text = 'Model_2 with Tanh activation and glorot Normal weight initialization'
3 diff_ops(openfromfile('history_22'),text)
```



from the above graph we can clearly observe that using a different activation function 'Tanh' and initializing the weights using glorot decreased the accuracy by 1%, though surprisingly the model performed better on test data than the train data after few initial epochs i.e generalization accuracy comes to be higher and loss lower in this case. Sudden increase in the loss around 3rd epoch can be understood from the concept of vanishing gradients.

```
In [38]: 1 text = 'Model_2 with RMSprop as optimizer'
        2 diff_ops(openfromfile('history_33'),text)
```



Here we see that using RMSProp as the optimizer instead of Adam has no significant impact, the accuracy on test is more or less same, also the execution time on test data is also same. As the data we are using is trivial, that's why RMSProp also converges fastly.

Conclusion

In [41]:

[illegible]

```

43 print('\n\n\n')
44
45
46
47

```

3 Different CNN architectures

				trainig loss	Training Accura
cy(%)	Test loss	Model	Test_Accuracy(%)		
1253	0.01964392685413427	2 Convolution layers + 1 pooling layer + 1 Dense layer + Adam	0.9951	0.009973527063508906	0.996700003147
3416	0.023159436680417637	5 Convolution layers + 2 pooling layer + 2 Dense Layers + Adam	0.9937	0.04901980113451524	0.986700010200
0813	0.01964392685413427	6 Convolution layers + 1 6 pooling layers + 2 Dense Layers + Adam	0.9951	0.0268287254908743	0.991850007375

Different output of best performing model

				trainig loss	Training
Accuracy(%)	Test loss	Model	Test_Accuracy(%)		
66678905487	0.03456393923192763	Model_2 without BatchNormalization and Dropout	0.9927	0.004473916407167261	0.99871
33438118299	0.03541449715318158	Model_2 with using Tanh activation and glorot normal weight initializer	0.9903	0.0761473451834172	0.97828
00112573306	0.046560718773310396	Model_2 with RMSProp optimizer	0.9907	0.09186855908615446	0.98205



#observation A 10 layer CNN which consists of 5 Convolutional layers and 2 maxpooling layers and Dense layers with relu activation and a final softmax output layer gives us the best accuracy of 99.5% on test data and a test loss of 0.018 .BatchNormalization and Dropout rate helps us in avoiding overfitting in such a complex architecture.This also gives us insight of how we can further tune the dropout rates kernel sizes and model with 7 Convolutional layers which is our 3rd model with.We also tried the GridSearch Cross validation on our first model with different kernel sizes,though it increases the training time by huge margin.A subsequent tuning may result in less time complexity

We also observe that removing BatchNormalization and dropout layers from the best performing model leads to overfitting and slow convergence by the model.using Tanh activation and glorot normal weight initialization decreases the accuracy by 1% to 98% where on an average we were getting the accuracy of 99%,which shos that how significant activation function is while initializing the model.RMSprop instead of AdamOptimizer does not have any major impact,and everything remains same more or less.